



Advanced SQL: Grouping, Aggregation and Views

Abdu Alawini

University of Illinois at Urbana-Champaign

CS411: Database Systems



Learning Objectives

After this lecture, you should be able to:

- Write grouping and aggregate SQL queries
- Define SQL *views*, differentiate between *virtual* and *materialized* views, and know when a view can be updated.



Aggregations

- SUM, AVG, COUNT, MIN, and MAX can be applied to a column in a SELECT clause to produce that aggregation on the column.
- Also, COUNT(*) counts the number of tuples.



Example: Aggregation

- From Sells(cafe, drink, price), find the average price of Mocha:

```
SELECT AVG(price)  
FROM Sells  
WHERE drink = 'Mocha';
```



Eliminating Duplicates in an Aggregation

- DISTINCT inside an aggregation causes duplicates to be eliminated before the aggregation.
- Example: find the number of different prices charged for Mocha:

```
SELECT COUNT(DISTINCT price)  
FROM Sells  
WHERE drink = 'Mocha';
```



NULL's Ignored in Aggregation

- NULL never contributes to a sum, average, or count, and can never be the minimum or maximum of a column.
- But if all the values in a column are null, then the result of the aggregation is NULL.



Example: Effect of NULL's

The number of cafes
that sell Mocha.

```
SELECT count(*)  
FROM Sells  
WHERE drink = 'Mocha';
```

```
SELECT count(price)  
FROM Sells  
WHERE drink = 'Mocha';
```

The number of cafes
that sell Mocha at a
known price.



Grouping

- We may follow a SELECT-FROM-WHERE expression by GROUP BY and a list of attributes.
- The relation that results from the SELECT-FROM-WHERE is grouped according to the values of all those attributes, and any aggregation is applied only within each group.

Example: Grouping

- From Sells(cafe, drink, price), find the average price for each drink:

```
SELECT drink, AVG(price)  
FROM Sells  
GROUP BY drink;
```



Example: Grouping

- From Sells(cafe, drink, price) and Frequents(customer, cafe), find for each customer the average price of Mocha at the cafes they frequent:



Example: Grouping

- From Sells(cafe, drink, price) and Frequents(customer, cafe), find for each customer the average price of Mocha at the cafes they frequent:

```
SELECT customer, AVG(price)  
FROM Frequents NATURAL JOIN Sells  
WHERE drink = 'Mocha'  
GROUP BY customer;
```

Compute combos
of frequents
and sells for
Mocha selling cafés,
then group
by customer.



Restriction on SELECT Lists With Aggregation

If any aggregation is used, then each element of the SELECT list must be either:

1. Aggregated, or
2. An attribute on the GROUP BY list.

Doesn't make sense to neither have be grouped or aggregated – how are you “collapsing” values down then?



HAVING Clauses

- HAVING <condition> may follow a GROUP BY clause.
- If so, the condition applies to each group, and groups not satisfying the condition are eliminated.



The HAVING clause: Example

```
SELECT drink, AVG(price)  
FROM Sells  
GROUP BY drink  
HAVING COUNT(cafe) >= 3
```

What does this query ask for?



The HAVING clause: Example

```
SELECT drink, AVG(price)  
FROM Sells  
GROUP BY drink  
HAVING COUNT(cafe) >= 3
```

What does this query ask for?

Average price of drinks that are sold at least at three cafes.



Requirements on HAVING Conditions

- These conditions may refer to any relation or tuple-variable in the FROM clause.
- They may refer to attributes of those relations, as long as the attribute makes sense within a group; i.e., it is either:
 1. A grouping attribute, or
 2. Aggregated.



General form of Grouping and Aggregation

```
SELECT      S  
FROM        R1,...,Rn  
WHERE       C1  
GROUP BY   a1,...,ak  
HAVING     C2
```

S = may contain attributes a_1, \dots, a_k and/or any aggregates but NO OTHER ATTRIBUTES

C_1 = is any condition on the attributes in R_1, \dots, R_n

C_2 = is any condition on aggregate expressions or grouping attributes



General form of Grouping and Aggregation

```
SELECT  S  
FROM    R1,...,Rn  
WHERE   C1  
GROUP BY a1,...,ak  
HAVING  C2
```

Evaluation steps:

1. Compute the FROM-WHERE part, obtain a table with all attributes in R_1, \dots, R_n
2. Group by the attributes a_1, \dots, a_k
3. Compute the aggregates in C_2 and keep only groups satisfying C_2
4. Compute aggregates in S and return the result



Outline

- ✓ Aggregation and Grouping
- Views



Views

- A view is a “virtual table,” a relation that is defined in terms of the contents of other tables and views.
- Declare by:

```
CREATE VIEW <name> AS <query>;
```

- Views are not stored in the database, but can be queried as if they existed.
 - We'll talk about an exception later
- In contrast, a relation whose value is really stored in the database is called a *base table*.



Example: View Definition

- CanDrink (customer, drink) is a view “containing” the customer-drink pairs such that the customer frequents at least one cafe that serves the drink on relations Frequent (customer, cafe) and Sells (cafe, drink, price)

```
CREATE VIEW CanDrink AS
    SELECT customer, drink
    FROM Frequent, Sells
    WHERE Frequent.cafe = Sells.cafe;
```



Example: Accessing a View

- You may query a view as if it were a base table.
- Example:

```
SELECT drink FROM CanDrink  
WHERE customer = 'Sally';
```



What's Useful about Views

1. Can be used as relations in other queries
 - Allows the user to query things that make more sense
 - Can be stored (materialized) as appropriate
 - Sometimes can even be updated!



What's Useful about Views

2. Can facilitate security/access control

- We can assign users permissions on different views
- Can select or project so we only reveal what we want!

3. Describe *transformations* or *mappings* from one schema (the base relations) to another (the output of the view)

- The basis of converting from different data models or representations
- Incredibly useful for logical data independence



View Example

A company's database includes a relation:

Part (PartID: Char(4), Weight:real)

- Weight is stored in Kilograms (kg)
- Company is purchased by a firm that uses imperial weights (lb)
- Databases must be integrated and use lb.
- But there's much old software using kilograms.
- 1 KG = ~2.2 LB

Example (Cont.)

Create a view over the parts table so that it uses imperial weights

```
CREATE VIEW Part_lb AS  
SELECT PartID, Weight*2.2 as Weight_lb  
FROM Part
```



Materialized Views

A **materialized view** is one that is computed once and its results are stored as a table

- Think of this as a cached answer
- These are incredibly useful!
- Techniques exist for using materialized views to answer other queries
- Materialized views are the basis of relating tables in different schemas

```
CREATE MATERIALIZED VIEW AS ...
```



Views Should Stay Fresh

- Views (sometimes called *intensional relations*) behave, from the perspective of a query language, exactly like base relations (*extensional relations*)
- But there's an association that should be maintained:
 - If tuples change in the base relation, they should change in the view (whether it's materialized or not)
 - If tuples change in the view, that should reflect in the base relation(s)



View Maintenance and Updates

- There exist algorithms to **incrementally** recompute a materialized view when the base relations change
- We can try to propagate view changes to the base relations
 - However, there are lots of views that aren't easily updatable:

| R | A | B |
|---|---|---|
| 1 | 1 | 2 |
| 2 | 2 | 2 |

| S | B | C |
|---|---|---|
| 2 | 2 | 4 |
| 2 | 2 | 3 |

| R \bowtie S | A | B | C |
|---------------|---|---|---|
| 1 | 1 | 2 | 4 |
| 1 | 1 | 2 | 3 |
| 2 | 2 | 2 | 4 |
| 2 | 2 | 2 | 3 |

delete?



The good news

We can ensure views are updatable by enforcing certain constraints:

- It is defined on a single base table
- Using only selection and projection
- No aggregates
- No DISTINCT

...but this limits the kinds of views we can have